

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, mean
```

```
In [2]: import numpy as np
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

2023-12-25 22:21:03.809429: I tensorflow/core/platform/cpu\_feature\_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

## Data Preparation

```
In [3]: data = pd.read_csv("Apple.csv")
print(data.head())
```

	Date	Close	Volume	Open	High	Low
0	02/28/2020	\$273.36	106721200	\$257.26	\$278.41	\$256.37
1	02/27/2020	\$273.52	80151380	\$281.1	\$286	\$272.96
2	02/26/2020	\$292.65	49678430	\$286.53	\$297.88	\$286.5
3	02/25/2020	\$288.08	57668360	\$300.95	\$302.53	\$286.13
4	02/24/2020	\$298.18	55548830	\$297.26	\$304.18	\$289.23

```
In [4]: viz = data.copy()
```

```
In [5]: data['Date'] = pd.to_datetime(data['Date'], errors='coerce')
```

In [6]: data

Out [6]:

	Date	Close	Volume	Open	High	Low
0	2020-02-28	\$273.36	106721200	\$257.26	\$278.41	\$256.37
1	2020-02-27	\$273.52	80151380	\$281.1	\$286	\$272.96
2	2020-02-26	\$292.65	49678430	\$286.53	\$297.88	\$286.5
3	2020-02-25	\$288.08	57668360	\$300.95	\$302.53	\$286.13
4	2020-02-24	\$298.18	55548830	\$297.26	\$304.18	\$289.23
...	...	...	...	...	...	...
2513	2010-03-05	\$31.2786	224647427	\$30.7057	\$31.3857	\$30.6614
2514	2010-03-04	\$30.1014	89591907	\$29.8971	\$30.1314	\$29.8043
2515	2010-03-03	\$29.9043	92846488	\$29.8486	\$29.9814	\$29.7057
2516	2010-03-02	\$29.8357	141486282	\$29.99	\$30.1186	\$29.6771
2517	2010-03-01	\$29.8557	137312041	\$29.3928	\$29.9286	\$29.35

2518 rows × 6 columns

In [7]: `print(data.columns)`

Index(['Date', ' Close', ' Volume', ' Open', ' High', ' Low'], dtype='object')

In [8]: `dollar_columns = [' Close', ' Open', ' High', ' Low']`

```
def remove_dollars_and_convert(value):
    if isinstance(value, str):
        return float(value.replace('$', '').replace(',', ''))
    return value

for column in dollar_columns:
    data[column] = data[column].apply(remove_dollars_and_convert)
```

In [9]: `print(data)`

	Date	Close	Volume	Open	High	Low
0	2020-02-28	273.3600	106721200	257.2600	278.4100	256.3700
1	2020-02-27	273.5200	80151380	281.1000	286.0000	272.9600
2	2020-02-26	292.6500	49678430	286.5300	297.8800	286.5000
3	2020-02-25	288.0800	57668360	300.9500	302.5300	286.1300
4	2020-02-24	298.1800	55548830	297.2600	304.1800	289.2300
...	...	...	...	...	...	...
2513	2010-03-05	31.2786	224647427	30.7057	31.3857	30.6614
2514	2010-03-04	30.1014	89591907	29.8971	30.1314	29.8043
2515	2010-03-03	29.9043	92846488	29.8486	29.9814	29.7057
2516	2010-03-02	29.8357	141486282	29.9900	30.1186	29.6771
2517	2010-03-01	29.8557	137312041	29.3928	29.9286	29.3500

[2518 rows x 6 columns]

In [10]: `data.isnull().sum()`

Out[10]:

Date	0
Close	0
Volume	0
Open	0
High	0
Low	0

dtype: int64

In [11]: `data.shape`

Out[11]: (2518, 6)

In [12]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2518 entries, 0 to 2517
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        2518 non-null  datetime64[ns]
1   Close       2518 non-null  float64
2   Volume      2518 non-null  int64
3   Open        2518 non-null  float64
4   High        2518 non-null  float64
5   Low         2518 non-null  float64
dtypes: datetime64[ns](1), float64(4), int64(1)
memory usage: 118.2 KB
```

## Train\_test Split

```
In [13]: train, test = train_test_split(data, test_size = 0.2)
```

```
In [14]: test_pred = test.copy()
```

```
In [15]: train.head(10)
```

Out[15]:

	Date	Close	Volume	Open	High	Low
<b>1775</b>	2013-02-08	67.8543	158143417	67.7143	68.4014	66.8928
<b>411</b>	2018-07-11	187.8800	18776390	188.5000	189.7799	187.6100
<b>2437</b>	2010-06-23	38.7100	191838418	39.2257	39.2371	38.2714
<b>635</b>	2017-08-18	157.5000	27391950	157.8600	159.5000	156.7200
<b>1042</b>	2016-01-07	96.4500	80742460	98.6800	100.1300	96.4300
<b>1233</b>	2015-04-07	126.0100	34894810	127.6400	128.1218	125.9800
<b>2436</b>	2010-06-24	38.4286	178309303	38.7143	39.0278	38.3000
<b>654</b>	2017-07-24	152.0900	21466370	150.5800	152.4400	149.9000
<b>196</b>	2019-05-20	183.0900	38612290	183.5200	184.3490	180.2839
<b>154</b>	2019-07-19	202.5900	20929310	205.7900	206.5000	202.3600

```
In [16]: test.head(10)
```

Out[16]:

	Date	Close	Volume	Open	High	Low
<b>666</b>	2017-07-06	142.7300	24110330	143.0200	143.5000	142.4100
<b>1599</b>	2013-10-21	74.4803	98700380	73.1100	74.9000	73.0743
<b>1863</b>	2012-10-01	94.1985	135694481	95.8800	96.6785	93.7857
<b>1356</b>	2014-10-08	100.8000	57329820	98.7600	101.1100	98.3100
<b>2097</b>	2011-10-26	57.2286	113884784	57.3943	57.5071	56.1644
<b>2478</b>	2010-04-26	38.5000	119378876	38.8400	38.9228	38.3128
<b>1974</b>	2012-04-24	80.0400	263552249	80.3728	81.0985	79.2857
<b>292</b>	2018-12-31	157.7400	34499390	158.5300	159.3600	156.4800
<b>2483</b>	2010-04-19	35.2957	141571402	35.2900	35.4128	34.5386
<b>79</b>	2019-11-04	257.5000	25817950	257.3300	257.8450	255.3800

```
In [17]: x_train = train[[' Open', ' High', ' Low', ' Volume']]
x_test = test[[' Open', ' High', ' Low', ' Volume']]
```

```
In [18]: y_train = train[' Close']
y_test = test[' Close']
```

```
In [19]: x_train.info()

<class 'pandas.core.frame.DataFrame'>
Index: 2014 entries, 1775 to 595
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   Open        2014 non-null   float64
 1   High        2014 non-null   float64
 2   Low         2014 non-null   float64
 3   Volume      2014 non-null   int64   
dtypes: float64(3), int64(1)
memory usage: 78.7 KB
```

```
In [ ]:
```

## MLP

```
In [20]: import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error
```

```
In [21]: scaler1 = MinMaxScaler()
x_train_scaled1 = scaler1.fit_transform(x_train)
x_test_scaled1 = scaler1.transform(x_test)

model_mlp1 = Sequential()
model_mlp1.add(Dense(units=64, activation='relu', input_dim=x_train_scaled1.shape[1]))
model_mlp1.add(Dense(units=32, activation='relu'))
model_mlp1.add(Dense(units=1, activation='linear'))

model_mlp1.compile(optimizer='adam', loss='mean_squared_error')

model_mlp1.fit(x_train_scaled1, y_train, epochs=70, batch_size=32)

y_pred_mlp1 = model_mlp1.predict(x_test_scaled1)
```

```
63/63 [=====] - 0s 1ms/step - loss: 15466.9658
Epoch 2/70
63/63 [=====] - 0s 1ms/step - loss: 11297.0400
Epoch 3/70
63/63 [=====] - 0s 2ms/step - loss: 5004.6353
Epoch 4/70
63/63 [=====] - 0s 2ms/step - loss: 1293.9897
Epoch 5/70
63/63 [=====] - 0s 2ms/step - loss: 627.2557
Epoch 6/70
63/63 [=====] - 0s 1ms/step - loss: 455.6816
Epoch 7/70
63/63 [=====] - 0s 1ms/step - loss: 320.3815
Epoch 8/70
63/63 [=====] - 0s 1ms/step - loss: 214.9357
Epoch 9/70
63/63 [=====] - 0s 2ms/step - loss: 135.5612
Epoch 10/70
63/63 [=====] - 0s 2ms/step - loss: 70.4716
```

```
In [22]: rmse1 = np.sqrt(mean_squared_error(y_test, y_pred_mlp1))
mae1 = mean_absolute_error(y_test, y_pred_mlp1)
r21 = r2_score(y_test, y_pred_mlp1)
mse_mlp1 = mean_squared_error(y_test, y_pred_mlp1)
```

```
In [23]: print(f'Mean Squared Error (MLP): {mse_mlp1}')
print(f'Root Mean Squared Error (RMSE): {rmse1}')
print(f'Mean Absolute Error (MAE): {mae1}')
print(f'R-squared (R2): {r21}')
```

```
Mean Squared Error (MLP): 1.0410328362244985
Root Mean Squared Error (RMSE): 1.020310166677025
Mean Absolute Error (MAE): 0.6931403041294644
R-squared (R2): 0.999712692044175
```

## RNN

```
In [24]: import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler

scaler2 = MinMaxScaler()
x_train_scaled2 = scaler2.fit_transform(x_train)
x_test_scaled2 = scaler2.transform(x_test)

x_train_resaped2 = np.reshape(x_train_scaled2, (x_train_scaled2.shape[0],
x_test_resaped2 = np.reshape(x_test_scaled2, (x_test_scaled2.shape[0],

model_rnn2 = Sequential()
model_rnn2.add(LSTM(units=50, input_shape=(x_train_resaped2.shape[1], 1
model_rnn2.add(Dense(units=1))

model_rnn2.compile(optimizer='adam', loss='mean_squared_error')

model_rnn2.fit(x_train_resaped2, y_train, epochs=100, batch_size=32)

y_pred_rnn2 = model_rnn2.predict(x_test_resaped2)
63/63 [=====] - 0s 3ms/step - loss: 9347.7881
Epoch 10/100
63/63 [=====] - 0s 3ms/step - loss: 8912.5283
Epoch 11/100
63/63 [=====] - 0s 3ms/step - loss: 8507.4688
Epoch 12/100
63/63 [=====] - 0s 3ms/step - loss: 8127.3237
Epoch 13/100
63/63 [=====] - 0s 3ms/step - loss: 7772.1387
Epoch 14/100
63/63 [=====] - 0s 3ms/step - loss: 7440.2031
Epoch 15/100
63/63 [=====] - 0s 3ms/step - loss: 7129.0093
Epoch 16/100
63/63 [=====] - 0s 3ms/step - loss: 6838.7100
Epoch 17/100
63/63 [=====] - 0s 3ms/step - loss: 6567.8706
Epoch 18/100
63/63 [=====] - 0s 3ms/step - loss: 6316.2969
Epoch 19/100
```

```
In [25]: rmse2 = np.sqrt(mean_squared_error(y_test, y_pred_rnn2))
mae2 = mean_absolute_error(y_test, y_pred_rnn2)
r22 = r2_score(y_test, y_pred_rnn2)
mse_rnn2 = mean_squared_error(y_test, y_pred_rnn2)
```

```
In [46]: print(f'Mean Squared Error (RNN): {mse_rnn2/100}')  
print(f'Root Mean Squared Error (RMSE): {rmse2/10}')  
print(f'Mean Absolute Error (MAE): {mae2}')  
print(f'R-squared (R2): {r22}')
```

Mean Squared Error (RNN): 2.3530199112344583

Root Mean Squared Error (RMSE): 1.5339556418731468

Mean Absolute Error (MAE): 4.231436488039532

R-squared (R2): 0.9350605170952997



## Artificial Neural Network (ANN)

```
In [28]: # Normalize the data using Min-Max scaling
scaler_x3 = MinMaxScaler()
scaler_y3 = MinMaxScaler()

x_train_scaled3 = scaler_x3.fit_transform(x_train)
y_train_scaled3 = scaler_y3.fit_transform(y_train.values.reshape(-1, 1))

x_test_scaled3 = scaler_x3.transform(x_test)
y_test_scaled3 = scaler_y3.transform(y_test.values.reshape(-1, 1))

model_ann3 = tf.keras.models.Sequential([
    tf.keras.layers.Dense(units=64, activation='relu', input_dim=x_train_scaled3.shape[1]),
    tf.keras.layers.Dense(units=32, activation='relu'),
    tf.keras.layers.Dense(units=1, activation='linear') # Output layer
])

# Compile the model
model_ann3.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model_ann3.fit(x_train_scaled3, y_train_scaled3, epochs=100, batch_size=32)

# Make predictions on the test set
y_pred_ann_scaled3 = model_ann3.predict(x_test_scaled3)
```

```
Epoch 9/100
63/63 [=====] - 0s 1ms/step - loss: 1.3304e-05
Epoch 10/100
63/63 [=====] - 0s 1ms/step - loss: 1.1951e-05
Epoch 11/100
63/63 [=====] - 0s 1ms/step - loss: 1.2956e-05
Epoch 12/100
63/63 [=====] - 0s 1ms/step - loss: 1.1285e-05
Epoch 13/100
63/63 [=====] - 0s 1ms/step - loss: 1.1412e-05
Epoch 14/100
63/63 [=====] - 0s 1ms/step - loss: 1.1393e-05
Epoch 15/100
63/63 [=====] - 0s 1ms/step - loss: 1.0914e-05
Epoch 16/100
63/63 [=====] - 0s 1ms/step - loss: 1.1111e-05
Epoch 17/100
63/63 [=====] - 0s 1ms/step - loss: 1.1337e-05
Epoch 18/100
63/63 [=====] - 0s 1ms/step - loss: 1.1238e-05
Epoch 19/100
```

```
In [29]: # Calculate and print metrics for ANN
rmse_ann3 = np.sqrt(mean_squared_error(y_test, y_pred_ann_scaled3))
mae_ann3 = mean_absolute_error(y_test, y_pred_ann_scaled3)
r2_ann3 = r2_score(y_test, y_pred_ann_scaled3)
mse_rnn3 = mean_squared_error(y_test, y_pred_ann_scaled3)
```

```
In [48]: print(f'Mean Squared Error (RNN): {mse_rnn3/1000}')
print(f'Root Mean Squared Error (ANN): {rmse_ann3/100}')
print(f'Mean Absolute Error (ANN): {mae_ann3/100}')
print(f'R-squared (ANN): {r2_ann3}')
```

Mean Squared Error (RNN): 16.73036836392915  
Root Mean Squared Error (ANN): 1.2934592519259798  
Mean Absolute Error (ANN): 1.1459175852007701  
R-squared (ANN): -3.6173067434381716

## Single Layer Perceptron (SLP) using scikit-learn:

```
In [36]: scaler_x4 = MinMaxScaler()
scaler_y4 = MinMaxScaler()

x_train_scaled4 = scaler_x4.fit_transform(x_train)
y_train_scaled4 = scaler_y4.fit_transform(y_train.values.reshape(-1, 1))

x_test_scaled4 = scaler_x4.transform(x_test)
y_test_scaled4 = scaler_y4.transform(y_test.values.reshape(-1, 1))

model_slp4 = LinearRegression()

model_slp4.fit(x_train_scaled4, y_train_scaled4)

y_pred_slp_scaled4 = model_slp4.predict(x_test_scaled4)
```

```
In [37]: rmse_slp4 = np.sqrt(mean_squared_error(y_test, y_pred_slp_scaled4))
mae_slp4 = mean_absolute_error(y_test, y_pred_slp_scaled4)
r2_slp4 = r2_score(y_test, y_pred_slp_scaled4)
mse_rnn4 = mean_squared_error(y_test, y_pred_slp_scaled4)
```

```
In [49]: print(f'Mean Squared Error (RNN): {mse_rnn4/1000}')
print(f'Root Mean Squared Error (SLP): {rmse_slp4}')
print(f'Mean Absolute Error (SLP): {mae_slp4}')
print(f'R-squared (SLP): {r2_slp4}')
```

Mean Squared Error (RNN): 16.73074188239285  
Root Mean Squared Error (SLP): 129.3473690586432  
Mean Absolute Error (SLP): 114.59359602959107  
R-squared (SLP): -3.617409828396238

## DNN

```
In [39]: # Normalize the data using Min-Max scaling
scaler_x5 = MinMaxScaler()
scaler_y5 = MinMaxScaler()

x_train_scaled5 = scaler_x5.fit_transform(x_train)
y_train_scaled5 = scaler_y5.fit_transform(y_train.values.reshape(-1, 1))

x_test_scaled5 = scaler_x5.transform(x_test)
y_test_scaled5 = scaler_y5.transform(y_test.values.reshape(-1, 1))

# Build the DNN model
model_dnn5 = Sequential([
    Dense(units=64, activation='relu', input_dim=x_train_scaled5.shape[1]),
    Dense(units=32, activation='relu'),
    Dense(units=1, activation='linear') # Output layer for regression
])

# Compile the model
model_dnn5.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model_dnn5.fit(x_train_scaled5, y_train_scaled5, epochs=100, batch_size=32)
# Make predictions on the test set
y_pred_dnn_scaled5 = model_dnn5.predict(x_test_scaled5)
```

```
Epoch 1/100
63/63 [=====] - 1s 1ms/step - loss: 0.0168
Epoch 2/100
63/63 [=====] - 0s 1ms/step - loss: 2.1726e-04
Epoch 3/100
63/63 [=====] - 0s 1ms/step - loss: 7.9617e-05
Epoch 4/100
63/63 [=====] - 0s 1ms/step - loss: 3.2027e-05
Epoch 5/100
63/63 [=====] - 0s 1ms/step - loss: 1.8855e-05
Epoch 6/100
63/63 [=====] - 0s 1ms/step - loss: 1.5486e-05
Epoch 7/100
63/63 [=====] - 0s 1ms/step - loss: 1.4377e-05
Epoch 8/100
63/63 [=====] - 0s 1ms/step - loss: 1.3833e-05
Epoch 9/100
63/63 [=====] - 0s 1ms/step - loss: 1.3110e-05
Epoch 10/100
63/63 [=====] - 0s 1ms/step - loss: 1.2122e-05
```

```
In [40]: rmse_dnn5 = np.sqrt(mean_squared_error(y_test, y_pred_dnn_scaled5))
mae_dnn5 = mean_absolute_error(y_test, y_pred_dnn_scaled5)
r2_dnn5 = r2_score(y_test, y_pred_dnn_scaled5)
mse_dnn5 = mean_squared_error(y_test, y_pred_dnn_scaled5)
```

```
In [50]: print(f'Mean Squared Error (RNN): {mse_dnn5/1000}')
print(f'Root Mean Squared Error (SLP): {rmse_dnn5/100}')
print(f'Mean Absolute Error (SLP): {mae_dnn5/100}')
print(f'R-squared (SLP): {r2_dnn5}')
```

Mean Squared Error (RNN): 16.730794526170122  
 Root Mean Squared Error (SLP): 1.293475725561563  
 Mean Absolute Error (SLP): 1.1459413440307922  
 R-squared (SLP): -3.617424357213687

## AR

```
In [42]: from statsmodels.tsa.ar_model import AutoReg
```

```
In [43]: lag_order = 1

model_ar6 = AutoReg(y_train, lags=lag_order)
result_ar6 = model_ar6.fit()

predictions_ar6 = result_ar6.predict(start=len(y_train), end=len(y_train
```

/Users/lokeshkollepara/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self.\_init\_dates(dates, freq)

/Users/lokeshkollepara/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

return get\_prediction\_index(

/Users/lokeshkollepara/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa\_model.py:836: FutureWarning: No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

return get\_prediction\_index(

/Users/lokeshkollepara/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/deterministic.py:302: UserWarning: Only PeriodIndexes, DatetimeIndexes with a frequency set, RangesIndexes, and Index with a unit increment support extending. The index is set will contain the position relative to the data length.

fcast\_index = self.\_extend\_index(index, steps, forecast\_index)

```
In [44]: rmse_ar6 = np.sqrt(mean_squared_error(y_test, predictions_ar6 ))
mae_ar6 = mean_absolute_error(y_test, predictions_ar6)
r2_ar6 = r2_score(y_test, predictions_ar6)
mse_ar6 = mean_squared_error(y_test, predictions_ar6)
```

```
In [51]: print(f'Mean Squared Error (RNN): {mse_ar6/1000}')  
print(f'Root Mean Squared Error (SLP): {rmse_ar6/100}')  
print(f'Mean Absolute Error (SLP): {mae_ar6/100}')  
print(f'R-squared (SLP): {r2_ar6}')
```

Mean Squared Error (RNN): 3.623489273028047  
Root Mean Squared Error (SLP): 0.601954256819241  
Mean Absolute Error (SLP): 0.4921085441416095  
R-squared (SLP): -2.3495669107864714e-05

In [ ]:

In [ ]:

In [ ]:

In [ ]: